

MINISTERIAL BRIEFING NOTE

SafeAI.sg Official Launch Dinner

Guest-of-Honour: Senior Minister of State Tan Kiat How

Date	Thursday, 21 May 2026
Time	6:45 PM – 9:30 PM
Venue	Restaurant Wakuda, Marina Bay Sands, Singapore
Host	SPARK by CIO Academy Asia, in partnership with Veeam
Convened by	David Chin, Chief Executive, SPARK by CIO Academy Asia
GOH Slot	7:20 PM – 7:35 PM (15 minutes, Guest-of-Honour Address)
Audience	~50 senior technology leaders: CIOs, CISOs, CTOs, Chief Risk Officers and AI practitioners from Singapore enterprise, government, and academia

1. PURPOSE OF THIS NOTE

SPARK by CIO Academy Asia respectfully requests confirmation of SMS Tan Kiat How's participation as Guest-of-Honour and is grateful for the office's continued support in elevating Singapore's industry-led AI governance agenda.

2. CONTEXT AND SIGNIFICANCE

What is SafeAI.sg?

SafeAI.sg is Singapore's industry-led Centre of Excellence for AI technical safety and practical governance, operationally anchored at the Punggol Digital District with institutional support from JTC Corporation, IMDA, and CSA. It is established and operated by SPARK by CIO Academy Asia — Singapore's senior technology leadership community representing approximately 16,000 CIOs, CTOs, and technology practitioners across the region.

SafeAI.sg bridges the gap between Singapore's national AI governance policy frameworks and the practitioner-level implementation realities faced by enterprise technology leaders deploying AI in production environments today. It is a non-regulatory, independent knowledge centre — providing risk-weighted, use-case-specific guidance, an open governance controls toolkit, a professional certification pathway, and a trusted community of practice.

Why this launch matters

The SafeAI.sg launch coincides with Singapore's position at the forefront of the global AI governance conversation. As a co-convenor of the International AI Safety Report — the largest global scientific collaboration on AI safety, backed by over 30 countries — Singapore has established itself as a credible voice in responsible AI deployment. SafeAI.sg translates that international standing into practical, on-the-ground infrastructure for Singapore's enterprise and public sector technology community.

3. SAFEAI.SG CHARTER OBJECTIVES — MINISTERIAL THREE-POINT BRIEF

SafeAI.sg delivers its mission through three mutually reinforcing delivery pillars that span the full lifecycle of AI governance readiness — from use-case risk assessment, through technology demonstration, to workforce capability building.

PILLAR 1 — Open Risk Assessment Toolkit and Governance Controls Standard

SafeAI.sg will develop and freely publish the Universal AI Use-Case Risk Score (UARS) — a self-service, five-dimension scoring instrument that any Singapore organisation can apply to any AI deployment to determine the appropriate tier of governance required. Alongside the UARS, SafeAI.sg publishes a full governance controls toolkit specifying commensurate controls, implementation

processes, and a policy-as-code extensibility framework, ensuring governance reaches every system the AI use case touches.

Outcome: Singapore enterprises and public agencies gain immediate, actionable, cost-free access to calibrated AI governance guidance — reducing both under-governance risk and innovation-stalling over-governance.

PILLAR 2 — Technology Showcase and Associate Programme

SafeAI.sg operates an Associate Programme for AI platform providers, cybersecurity vendors, observability tooling companies, and governance technology firms to demonstrate how strong AI governance controls can be embedded into real-world enterprise deployments. Associates showcase their technologies through community events, reference architecture demonstrations, and case study publications — illustrating how their products address specific governance requirements at defined risk tiers.

Outcome: Singapore becomes a credible regional demonstration environment for enterprise AI governance — supporting EDB, JTC, and IMDA's positioning of Singapore as the regional hub for trusted, responsibly deployed AI.

PILLAR 3 — Professional Certification and Skills Development

SafeAI.sg will co-develop a structured professional certification programme in partnership with Singapore's institutes of higher learning, spanning four levels — from career conversion for non-technical professionals through to senior certification for enterprise architects and AI governance leaders. Governance frameworks are only sustainable when the people responsible for implementing them possess the knowledge to do so.

Outcome: Singapore builds a sustainable, skilled AI governance workforce — ensuring that responsible AI deployment is an organisational capability, not a compliance exercise dependent on specialist consultants.

4. PROGRAMME OUTLINE — 21 MAY 2026

Time	Duration	Item
6:45 PM	30 min	Registration & Networking
7:15 PM	5 min	Opening Remarks — David Chin, Chief Executive, SPARK by CIO Academy Asia; Casa Goh, Country Manager, Veeam Singapore
7:20 PM	15 min	Guest-of-Honour Address — Senior Minister of State Tan Kiat How, Ministry of Digital Development and Information & Ministry of Health ★ GOH SLOT
7:35 PM	20 min	Special Keynote / Fireside Chat — Stephen Clare, Lead Writer, International AI Safety Report (2025/2026); Affiliate, Centre for the Governance of AI (GovAI)
7:55 PM	30 min	Keynote — Veeam
8:25 PM	30 min	Practitioner Panel — SafeAI.sg Advisory Council & Tech Advisory Panel Members (moderated by David Chin)
8:55 PM	10 min	Acknowledgement & Presentation of Plaques — SafeAI.sg Advisory Council and Tech Advisory Panel members (Presenter: Council Chair)
9:05 PM	40 min	Networking & Close

5. FEATURED SPEAKERS

Guest-of-Honour

Senior Minister of State Tan Kiat How

Ministry of Digital Development and Information & Ministry of Health

SMS Tan Kiat How will deliver the Guest-of-Honour address at the official launch of SafeAI.sg. His remarks will lend ministerial endorsement to Singapore's industry-led AI governance agenda and contextualise SafeAI.sg's work within Singapore's broader national AI strategy and international commitments to responsible AI deployment.

Special Fireside Chat

Stephen Clare

Lead Writer, International AI Safety Report (2025/2026) | Affiliate, Centre for the Governance of AI (GovAI)

Stephen Clare is one of the lead architects of the International AI Safety Report — the largest global scientific collaboration on AI safety, backed by more than 30 countries and international organisations. He has synthesised insights from over 100 AI safety experts to produce a landmark assessment of the risks posed by general-purpose AI. His appearance at the SafeAI.sg launch directly connects Singapore's practitioner community to the frontline of global AI safety science, and reinforces Singapore's standing as a co-convenor of the international AI safety agenda.

Practitioner Panel

The practitioner panel brings together members of the SafeAI.sg Advisory Council and Tech Advisory Panel to discuss the building blocks of enterprise AI governance — spanning security, data trust, and agentic AI readiness. Confirmed panellists include:

Speaker	Organisation	Role
Dr. Adrian Tang	Centre for Strategic Infocomm Technologies (CSIT)	Group Director, Digital Defence Hub Advisory Council Member (in-principle)
Nikhil Dwarakanath	Grab	Group Head of Data & Analytics Advisory Council Member (in-principle)
Cheri Lim	Temasek	Chief Information Security Officer Advisory Council Member (confirmed)
Glen Francis	Google	CTO for Cloud, Asia Pacific Tech Advisory Panel Member (confirmed)
Jenny Koh	Snowflake Singapore	Country Manager Tech Advisory Panel Member (confirmed)

6. ACKNOWLEDGEMENT PLAQUES

At the conclusion of the formal programme, SPARK by CIO Academy Asia will formally acknowledge and present commemorative plaques to all SafeAI.sg Advisory Council members and Tech Advisory Panel members in attendance. This is a recognition of their commitment to Singapore's AI governance ecosystem as founding contributors to SafeAI.sg.

The presentation will be made by the SafeAI.sg Council Chair and takes approximately 10 minutes. SMS Tan Kiat How is warmly invited to remain for the presentation as a symbolic gesture of ministerial recognition for the Council's voluntary contribution to this national initiative.

7. STAKEHOLDER ENGAGEMENT — SENIOR LEADERSHIP BRIEFINGS COMPLETED

In the lead-up to the SafeAI.sg launch, David Chin has personally briefed the senior management teams of the following government agencies, private sector organisations, and technology partners. These briefings reflect the breadth of institutional buy-in that SafeAI.sg has secured across Singapore's public and private sector technology leadership community.

Government Agencies

Organisation	Senior Leader Briefed	Role
IMDA / DiSG	DCE Kiren Kumar	Deputy Chief Executive, Digital Industry Singapore
CSA	BG Edward Chen	Former Deputy CE (National Cyber Resilience), CSA; transitioning to Chief AI Officer, NCS
JTC Corporation	ACE David Tan	Assistant Chief Executive Officer
CSIT	Dr. Adrian Tang	Director, Digital Defence Hub

Private Sector Organisations

Organisation	Senior Leader Briefed	Role
Temasek	Cheri Lim	Chief Information Security Officer
Temasek	Baskara Rao	Chief Technology Officer
DBS Bank	Luis Carlos Cruz Huertas	Managing Director, Group Head of Data & AI Platform; Chief AI Officer
Keppel Technologies	Andrew Cheng	Chief Executive Officer
Keppel	Jason Chin	Group Chief Information Officer
Grab	Nikhil Dwarakanath	Group Head of Data & AI

Technology Partners

Organisation	Senior Leader Briefed	Role
Veeam	Beni Sia	Senior Vice President, Asia Pacific & Japan
Google	Glen Francis	CTO for Cloud, Asia Pacific
Snowflake Singapore	Jenny Koh	Country Manager
F5	Michael Quek	Regional Vice President
Fortinet	Jess Ng	Country Manager, Fortinet Singapore

8. SUGGESTED TALKING POINTS FOR SMS TAN KIAT HOW

The following five points are offered as a guide for SMS Tan's remarks. They may be adapted freely for tone and delivery.

AI Safety Is a Foundation, Not an Afterthought

1

Singapore's enterprises and public agencies are accelerating AI adoption — moving from pilots to production, from experimentation to enterprise-wide deployment. The risk is not in moving fast. The risk is in treating AI governance as something to retrofit once problems emerge. SafeAI.sg exists to change that default: every organisation embarking on an AI journey — regardless of size, sector, or stage — should establish its governance foundation at the outset. The cost of building AI safely from the start is a fraction of the cost of remediating ungoverned AI systems after the fact. Singapore's long-term competitiveness in AI depends not just on how fast we adopt AI, but on how responsibly we deploy it.

Suggested emphasis: Encourage every CIO, CISO, and technology leader in the room to treat AI safety as a design principle — built in from Day One of every AI initiative, not bolted on at the end.

The AI Landscape Evolves Faster Than Any Single Organisation Can Track — Collaboration Is the Only Sustainable Answer

2

AI risk is not static. The threat landscape, the capability landscape, and the regulatory landscape are all evolving simultaneously — faster than any individual organisation, regulator, or government agency can monitor alone. The 'unknown unknowns' of AI risk — failure modes we have not yet encountered, attack vectors that do not yet have names, systemic dependencies we have not yet mapped — represent the most significant governance challenge of our time. SafeAI.sg is built on the conviction that the only sustainable response is meaningful, structured collaboration: between industry practitioners who observe

AI systems in the field every day, and the regulators and policy agencies responsible for the frameworks that govern them. Singapore has the institutional relationships and the trust ecosystem to make this work. SafeAI.sg provides the structured channel for that exchange to happen continuously — not just at point-in-time consultations.

Suggested emphasis: *Call on regulators and industry to treat each other as essential intelligence partners — and commit to using SafeAI.sg as a structured, ongoing feedback channel to stay ahead of emerging AI risks together.*

Governance Frameworks Only Have Value If They Can Be Operationalised

3

Singapore has invested significantly in building thoughtful AI governance frameworks — the IMDA Model AI Governance Framework, the CSA Guidelines on Securing AI Systems, and our broader national AI strategy. These are well-considered, comprehensive foundations. The challenge we now face is the last mile: translating written governance principles into implemented technical controls, deployed consistently and verifiably across every system that AI touches. SafeAI.sg closes that gap. Its open toolkit — the Universal AI Use-Case Risk Score and accompanying governance controls standard — gives every technology team a practical, self-service instrument to implement the right controls for the right use case. It bridges the space between policy and practice, turning guidance into governance that is actually running in production.

Suggested emphasis: *Highlight that SafeAI.sg does not replace or compete with existing national frameworks — it operationalises them, giving the practitioner community the tools to implement what the frameworks recommend, in the specific context of their actual AI deployments.*

Risk-Based Governance Is the Only Approach That Scales

4

Not all AI carries the same risk. A customer-facing chatbot that answers FAQs carries very different risk from an autonomous agent approving financial transactions or routing clinical referrals. Governing both with the same intensity — applying enterprise-grade safety controls uniformly across every AI use case — is neither necessary nor sustainable. It creates unnecessary compliance cost, stalls innovation, and drives AI initiatives underground where they are governed even less. SafeAI.sg's risk-based approach is calibrated to this reality. The UARS framework enables organisations to apply proportionate governance: lightweight controls for low-risk applications, rigorous safety stacks for high-stakes deployments. The result is a governance model organisations can sustain at scale — not a compliance burden that makes responsible AI economically unviable and therefore not actually adopted.

Suggested emphasis: *Frame risk-based governance not as a compromise on safety, but as the only approach that keeps safety economically rational — and therefore actually implemented everywhere, consistently, at scale across Singapore's enterprise ecosystem.*

Industry-Led, Government-Anchored: Singapore's Model for Trusted AI

5

SafeAI.sg represents a distinctly Singaporean model for AI governance: industry-led, with deep government partnership, anchored in the practitioner community that actually deploys AI systems every day. The Advisory Council assembled tonight — spanning CSA, IMDA, JTC, CSIT, MDDI, AI Singapore, the Singapore AI Safety Institute, Temasek, DBS, Grab, and Google — is not a symbolic gesture. It is a working governance body, stress-testing our frameworks against real enterprise and public sector deployments, and providing structured feedback to the institutions that develop Singapore's AI guidelines. This model — practitioners and policymakers in the same room, working on the same problems — is Singapore's competitive advantage in AI governance. It is how Singapore ensures its frameworks keep pace with the technology, and how Singapore's enterprises can deploy AI with the confidence that the guardrails have been tested by people who have actually built these systems.

Suggested emphasis: *Close with a call to action: to every leader in the room — this is your community. The knowledge you carry from your own AI deployments, the failures as much as the successes, is exactly what makes SafeAI.sg work. Contribute to it. Singapore's AI future is built by the people in this room.*

9. LOGISTICS & PROTOCOL NOTES

- Dress Code: Smart Business / Business Formal

- Venue: Restaurant Wakuda, Marina Bay Sands, Singapore — private dining room reserved for the SafeAI.sg Launch Dinner
- Arrival: The GOH and team are invited to arrive by 7:00 PM for a brief pre-event meet with the host, David Chin (CEO, SPARK by CIO Academy Asia)
- Photography: A media photographer will be present. The GOH's team should advise on any specific photography protocols
- Accompaniment: SPARK will designate a host liaison to meet and accompany the GOH's team from the venue entrance
- Speech Notes: Suggested speaking notes and a two-page background brief on SafeAI.sg can be provided to the GOH's speechwriter upon request

10. ACTION REQUESTED

SPARK by CIO Academy Asia respectfully requests the following from SMS Tan Kiat How's office:

#	Action	By
1	Confirm SMS Tan Kiat How's participation as Guest-of-Honour and GOH speaker for the SafeAI.sg Launch Dinner on 21 May 2026	SMS's Office
2	Advise on any protocol, security, or photography requirements for the GOH's attendance	SMS's Office
3	Confirm if the GOH's team requires background speaking notes or a SafeAI.sg charter brief for speechwriting purposes	SMS's Office
4	Confirm attendance of Ministry representative for Advisory Council plaque presentation	SMS's Office

Contact for Responses and Coordination

David Chin | Chief Executive, SPARK | Founder, SafeAI.sg

Sheri Tan | Chief of Staff, SPARK | Chief Secretariat, SafeAI.sg

This briefing note is prepared in confidence for the office of Senior Minister of State Tan Kiat How. Please direct any queries to the contact above.